




Mediatorless Cross-DB Query Execution

Jan Hanack | BDAPRO WiSe 21/22 | 24.02.2022



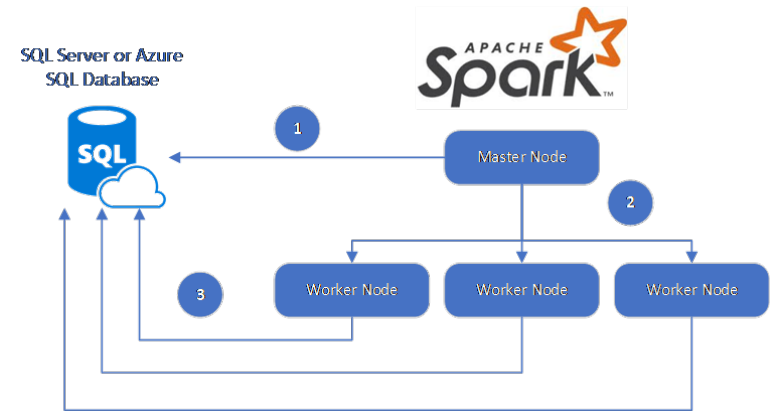
Recap & Relevance

- Use of Data Federations
 - Inter-Organizational
 - Intra-Organizational
- Enrich data analysis with external data
 - e.g. weather data for insurers, financial data, climate data for research
- Data Protection while cooperating
- Globalization & Big Data




Problem Statement

- Currently Mediator-Wrapper approach used
- Additional Data Transfer necessary
 - Increasing overhead and latency
 - Complicates Data Compliance
 - Bad fit for large scope, decentralized collaboration systems
- Akka's Actor Model non-ideal solution



Goals

- Develop Mediator-less approach
- As little as possible data copied & transferred
- Interoperability (e.g PostgreSQL, MariaDB)
- Similar to AWS Data Exchange or Snowflake
- Use output of Agora's execution planner

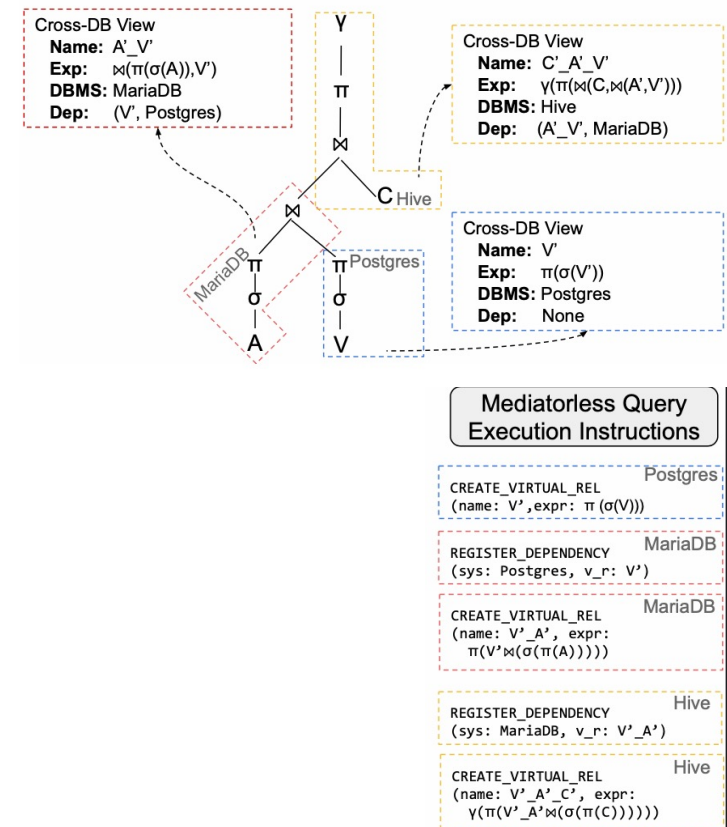


AGORA

— A Data Ecosystem for AI Innovation —

Solution & Contributions - Mediatorless Cross-DB Query Execution


- Interoperability between different DBMS
- Distributed setting possible
- DBMS and nodes communicate directly
- Foreign Tables
 - Point to Table in different DB
- Views
 - Generate Sub-Query Result






Solution & Contributions – Inter-Node/DB-Communication

- Protocol Buffers
 - Serializes structured Data
 - Increases efficiency for transfer
- gRPC
 - Uses Protocol Buffers
 - Client & Server Bindings
 - Easy Specification of services
 - Communication between Clients (DB Nodes)
 - Allows for millions of RPCs per second



Solution & Contributions – Architecture

- Slim server implementation per node
- One Client per Node and user
- User calls `createForeignTable()` and `createView()`
- DBs communicate directly
- Final Query result sent to user
- One or many DBs per node



Solution & Contributions – XDBX Implementation

- DBManager maps
 - DB ID → e.g. localhost:8080 → DB Type (e.g. POSTGRES)
 - One connector class per DB brand
 - getColumnns() converts to standard format and back
 - Two connections maintained:
 1. User ↔ Node Client

```
new SimpleXDBConnection( address: "localhost", port: 7000, username: "max", authentication: "123");
```

2. Client ↔ Database

- JDBC or ODBC →

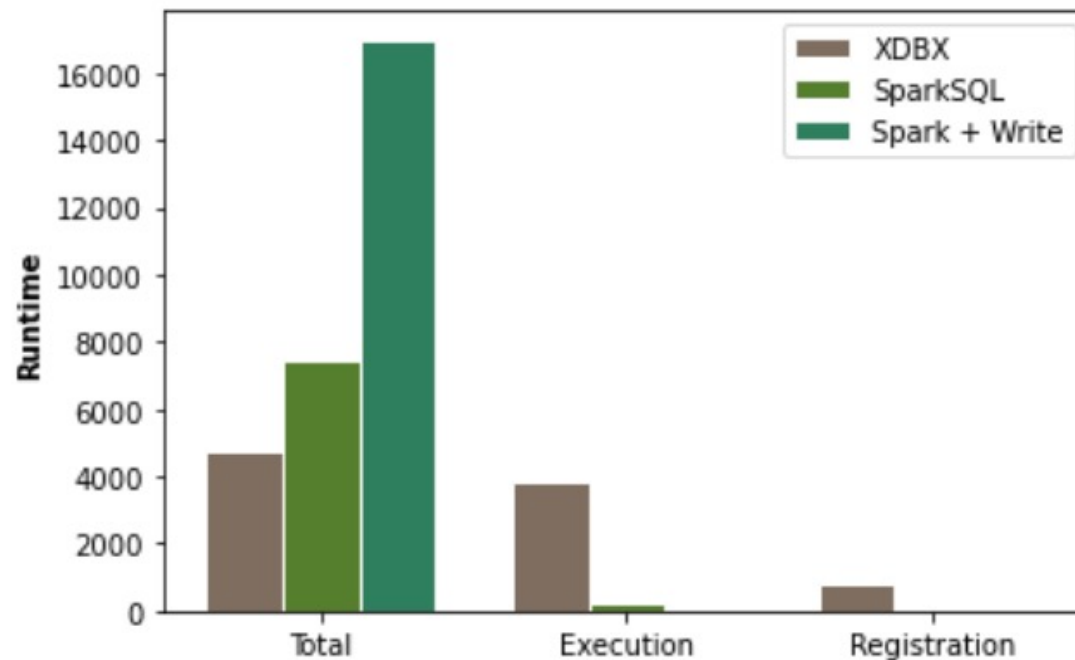
```
new PostgresInstance(clientFactory,  
  dbAddress: "localhost",  
  dbPort: 8081,  
  dbUsername: "bdapro_user",  
  dbPassword: "bdapro_password",  
  database: "bdapro_database");
```




Solution & Contributions – Experiments Setup

- Two virtual nodes
- 2-3 PostgreSQL containers via JDBC
- TPC-H Data & Queries
 - Common Benchmarking Tool
 - Scale Factor 1
 - Data distributed over DBs
- Benchmarks
 - Apache SparkSQL
 - Presto
 - Apache Calcite

Solution & Contributions – Benchmarks - TPC-H Query 3



- Machine: 6 Core CPU, 16 GB RAM



Solution & Contributions – Summary & Evaluation

- System works, query output correct
- Further improvement by use of indices
- For Report:
 - More TPC-H queries
 - Benchmark scale factor 10
 - Different DB setup
- Performance depends on workload (filters, or data manipulation, disk writes)
- Spark can be used on top of XDBX



Thank You for your Attention!